

Master Thesis

An LLVM Backend for Accelerate



Programming Languages and Compiler Construction
Department of Computer Science
Christian-Albrechts-University of Kiel

student: **Timo von Holtz**
advised by: Priv.-Doz. Dr. Frank Huch
Assoc. Prof. Dr. Manuel M. T. Chakravarty

Kiel, April 17, 2014

Selbstständigkeitserklärung

Ich erkläre hiermit, dass ich die vorliegende Arbeit selbstständig und nur unter Verwendung der angegebenen Literatur und Hilfsmittel angefertigt habe.

April 17, 2014

Timo von Holtz

Todo list

Write Introduction	1
------------------------------	---

Contents

1. Introduction	1
2. Technologies	3
2.1. LLVM	3
2.1.1. LLVM IR	3
2.1.2. Optimization	4
2.1.3. Vectorization	4
2.2. Accelerate	4
3. Contributions	5
3.1. llvm-general-quote	5
4. Conclusion	9
4.1. Related Work	9
A. Listings	11
Bibliography	15

List of Figures

2.1. sum as a C function	3
2.2. sum as a LLVM	4
3.1. Monadic generation of for loop	6
3.2. For Loop using <i>llvm-general-quote</i>	6
3.3. Expanded For Loop	7

List of Tables

1. Introduction

Write Introduction

2. Technologies

2.1. LLVM

LLVM[Lat02] is a compiler infrastructure written in C++. In contrast to GCC it is designed to be used as a library by compilers. Originally implemented for C and C++, the language-agnostic design (and the success) of LLVM has since spawned a wide variety of front ends: languages with compilers that use LLVM include ActionScript, Ada, D, Fortran, OpenGL Shading Language, Haskell, Java bytecode, Julia, Objective-C, Python, Ruby, Rust, Scala and C.

2.1.1. LLVM IR

LLVM defines its own language to represent programs. It uses Static Single Assignment (SSA) form.[AWZ88; RWZ88] A program is said to be in SSA form if each of its variables is defined exactly once, and each use of a variable is dominated by that variable's definition. SSA form greatly simplifies many dataflow optimizations because only a single definition can reach a particular use of a value, and finding that definition is trivial.

To get idea of how this looks like in practice, let's look at an example. Figure 2.1 shows a simple C function to sum up the elements of an array. The corresponding LLVM code is shown in figure 2.2.

```
double dotp(double* a, double* b, int length) {
    double x = 0;
    for (int i=0;i<length;i++) {
        x += a[i]*b[i];
    }
    return x;
}
```

Figure 2.1.: sum as a C function

```

define double @sum(double* %a, i32 %length) {
    %1 = icmp sgt i32 %length, 0
    br i1 %1, label %.lr.ph, label %._crit_edge

.lr.ph:
    ; preds = %0, %.lr.ph
    %indvars.iv = phi i64 [ %indvars.iv.next, %.lr.ph ], [ 0, %0 ]
    %x.01 = phi double [ %4, %.lr.ph ], [ 0.000000e+00, %0 ]
    %2 = getelementptr double* %a, i64 %indvars.iv
    %3 = load double* %2
    %4 = fadd double %x.01, %3
    %indvars.iv.next = add i64 %indvars.iv, 1
    %lftr.wideiv = trunc i64 %indvars.iv.next to i32
    %exitcond = icmp eq i32 %lftr.wideiv, %length
    br i1 %exitcond, label %._crit_edge, label %.lr.ph

._crit_edge:
    ; preds = %.lr.ph, %0
    %x.0.lcssa = phi double [ 0.000000e+00, %0 ], [ %4, %.lr.ph ]
    ret double %x.0.lcssa
}

```

Figure 2.2.: sum as a LLVM

2.1.2. Optimization

2.1.3. Vectorization

2.2. Accelerate

3. Contributions

3.1. *llvm-general-quote*

When writing a compiler using LLVM in Haskell there is a good tutorial on how to do it at <http://www.stephendiehl.com/llvm/>. It uses *llvm-general* to interface with LLVM. The general idea is to use a monadic generator to produce the AST on the fly.

Figure 3.1 shows how to implement a simple for loop using monadic generators. As you can tell this is much boilerplate code. We have to define the basic blocks manually and add the instructions one by one. This has some obvious drawbacks, as the code can get unreadable pretty quickly.

A solution is to use quasiquotation[Mai07] instead. The idea behind quasiquotation is, that you can define a DSL with arbitrary syntax, which you can then directly transform into Haskell data structures. This is done at compile-time, so you get the same type safety as writing the AST by hand.

I implemented *llvm-general-quote*, a quasiquotation library for LLVM. Figure 3.2 shows a for loop using my library.

Figure 3.3 shows the resulting LLVM IR. This is clearly more readable. Furthermore, one can see much more clearly what the produced code will be.

Another advantage of quasiquotation is antiquotation. This means you can still reference arbitrary Haskell variables from within the quotation. Using this the following are equivalent:

- `[llinstr| add i64 %x, 1 |]`
- `let y = 1 in [llinstr| add i64 %x, $opr:(y) |]`

The design of *llvm-general-quote* is inspired by *language-c-quote*, which is also used in the cuda implementation of Accelerate. I use “Happy” and *Alex*.

```

for :: Type                                -- type of the index
    → Operand                                -- starting index
    → (Operand → CodeGen Operand)           -- loop test to keep going
    → (Operand → CodeGen Operand)           -- increment the index
    → (Operand → CodeGen ())                -- body of the loop
    → CodeGen ()

for ti start test incr body = do
    loop ← newBlock "for.top"
    exit ← newBlock "for.exit"

    -- entry test
    c ← test start
    top ← cbr c loop exit

    -- Main loop
    setBlock loop
    c_i ← freshName
    let i = local c_i
    body i
    i' ← incr i
    c' ← test i'
    bot ← cbr c' loop exit
    – ← phi loop c_i ti [(i', bot), (start, top)]
    setBlock exit

```

Figure 3.1.: Monadic generation of for loop

```

[llgl]
define i64 @foo(i64 %start, i64 %end) {
  entry:
    br label %for

for:
  for i64 %i in %start to %end with i64 [0,%entry] as %x {
    %y = add i64 %i, %x
    ret i64 %y
  }
}
[]

```

Figure 3.2.: For Loop using *llvm-general-quote*

```

define i64 @foo(i64 %start, i64 %end) {
entry:
    br label %for

for:
    ; preds = %for.body, %entry
    %i = phi i64 [ %i.new, %for.body ], [ %start, %entry ]
    %x = phi i64 [ %y, %for.body ], [ 0, %entry ]
    %for.cond = icmp ule i64 %i, %end
    %i.new = add nuw nsw i64 %i, 1
    br i1 %for.cond, label %for.body, label %for.end

for.end:
    ; preds = %for
    ret i64 %x

for.body:
    ; preds = %for
    %y = add i64 %i, %x
    br label %for
}

```

Figure 3.3.: Expanded For Loop

4. Conclusion

4.1. Related Work

A. Listings

Listings

```
1 ; Function Attrs: nounwind readonly
2 define double @sum(double* nocapture readonly %a, i32 %length) #0 {
3     %l = icmp sgt i32 %length, 0
4     br i1 %l, label %lr.ph.preheader, label %._crit_edge
5
6     .lr.ph.preheader:                                     ; preds = %0
7     %2 = add i32 %length, -1
8     %3 = zext i32 %2 to i64
9     %4 = add i64 %3, 1
10    %end.idx = add i64 %3, 1
11    %n.vec = and i64 %4, 8589934576
12    %cmp.zero = icmp eq i64 %n.vec, 0
13    br i1 %cmp.zero, label %middle.block, label %vector.body
14
15    vector.body:                                          ; preds = %
16        .lr.ph.preheader, %vector.body
17    %index = phi i64 [ %index.next, %vector.body ], [ 0, %lr.ph.preheader ]
18    %vec.phi = phi <4 x double> [ %13, %vector.body ],
19        [ zeroinitializer, %lr.ph.preheader ]
20    %vec.phi6 = phi <4 x double> [ %14, %vector.body ],
21        [ zeroinitializer, %lr.ph.preheader ]
22    %vec.phi7 = phi <4 x double> [ %15, %vector.body ],
23        [ zeroinitializer, %lr.ph.preheader ]
24    %vec.phi8 = phi <4 x double> [ %16, %vector.body ],
25        [ zeroinitializer, %lr.ph.preheader ]
26    %5 = getelementptr double* %a, i64 %index
27    %6 = bitcast double* %5 to <4 x double>*
28    %wide.load = load <4 x double>* %6, align 8
29    %sum22 = or i64 %index, 4
30    %7 = getelementptr double* %a, i64 %sum22
31    %8 = bitcast double* %7 to <4 x double>*
32    %wide.load9 = load <4 x double>* %8, align 8
33    %sum23 = or i64 %index, 8
34    %9 = getelementptr double* %a, i64 %sum23
35    %10 = bitcast double* %9 to <4 x double>*
36    %wide.load10 = load <4 x double>* %10, align 8
37    %sum24 = or i64 %index, 12
38    %11 = getelementptr double* %a, i64 %sum24
39    %12 = bitcast double* %11 to <4 x double>*
40    %wide.load11 = load <4 x double>* %12, align 8
41    %13 = fadd <4 x double> %vec.phi, %wide.load
42    %14 = fadd <4 x double> %vec.phi6, %wide.load9
43    %15 = fadd <4 x double> %vec.phi7, %wide.load10
44    %16 = fadd <4 x double> %vec.phi8, %wide.load11
```

```

44 %index.next = add i64 %index, 16
45 %17 = icmp eq i64 %index.next, %n.vec
46 br i1 %17, label %middle.block, label %vector.body, !llvm.loop !0
47
48 middle.block:                                ; preds = %vector.body, %
      .lr.ph.preheader
49 %resume.val = phi i64 [ 0, %lr.ph.preheader ], [ %n.vec, %vector.body ]
50 %rdx.vec.exit.phi = phi <4 x double> [ zeroinitializer, %lr.ph.preheader
      ],
51                                     [ %13, %vector.body ]
52 %rdx.vec.exit.phi14 = phi <4 x double> [ zeroinitializer, %
      .lr.ph.preheader ],
53                                     [ %14, %vector.body ]
54 %rdx.vec.exit.phi15 = phi <4 x double> [ zeroinitializer, %
      .lr.ph.preheader ],
55                                     [ %15, %vector.body ]
56 %rdx.vec.exit.phi16 = phi <4 x double> [ zeroinitializer, %
      .lr.ph.preheader ],
57                                     [ %16, %vector.body ]
58 %bin.rdx = fadd <4 x double> %rdx.vec.exit.phi14, %rdx.vec.exit.phi
59 %bin.rdx17 = fadd <4 x double> %rdx.vec.exit.phi15, %bin.rdx
60 %bin.rdx18 = fadd <4 x double> %rdx.vec.exit.phi16, %bin.rdx17
61 %rdx.shuf = shufflevector <4 x double> %bin.rdx18, <4 x double> undef, <4
      x i32> <i32 2, i32 3, i32 undef, i32 undef>
62 %bin.rdx19 = fadd <4 x double> %bin.rdx18, %rdx.shuf
63 %rdx.shuf20 = shufflevector <4 x double> %bin.rdx19, <4 x double> undef,
      <4 x i32> <i32 1, i32 undef, i32 undef, i32 undef>
64 %bin.rdx21 = fadd <4 x double> %bin.rdx19, %rdx.shuf20
65 %18 = extractelement <4 x double> %bin.rdx21, i32 0
66 %cmp.n = icmp eq i64 %end.idx, %resume.val
67 br i1 %cmp.n, label %._crit_edge, label %lr.ph
68
69 .lr.ph:                                ; preds = %middle.block,
      %lr.ph
70 %indvars.iv = phi i64 [ %indvars.iv.next, %lr.ph ], [ %resume.val, %
      middle.block ]
71 %x.01 = phi double [ %21, %lr.ph ], [ %18, %middle.block ]
72 %19 = getelementptr double@%a, i64 %indvars.iv
73 %20 = load double@%19, align 8
74 %21 = fadd fast double %x.01, %20
75 %indvars.iv.next = add i64 %indvars.iv, 1
76 %lftr.wideiv1 = trunc i64 %indvars.iv.next to i32
77 %exitcond2 = icmp eq i32 %lftr.wideiv1, %length
78 br i1 %exitcond2, label %._crit_edge, label %lr.ph, !llvm.loop !3
79
80 ._crit_edge:                                ; preds = %lr.ph, %
      middle.block, %0
81 %x.0.lcssa = phi double [ 0.000000e+00, %0 ], [ %21, %lr.ph ], [ %18, %
      middle.block ]
82 ret double %x.0.lcssa
83 }
84
85 attributes #0 = { nounwind readonly }
86

```

```
87 | !0 = metadata !{metadata !0, metadata !1, metadata !2}
88 | !1 = metadata !{metadata !"llvm.vectorizer.width", i32 1}
89 | !2 = metadata !{metadata !"llvm.vectorizer.unroll", i32 1}
90 | !3 = metadata !{metadata !3, metadata !1, metadata !2}
```


Bibliography

- [AWZ88] Bowen Alpern, Mark N Wegman, and F Kenneth Zadeck. “Detecting equality of variables in programs”. In: *Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM. 1988, pp. 1–11. URL: <https://courses.cs.washington.edu/courses/cse501/04wi/papers/alpern-popl88.pdf>.
- [Die14] Stephen Diehl. *Implementing a JIT Compiled Language with Haskell and LLVM*. Jan. 2014. URL: <http://www.stephendiehl.com/llvm/>.
- [DM] Chris Dornan and Simon Marlow. *Alex: A lexical analyser generator for Haskell*. URL: <http://www.haskell.org/alex/>.
- [GM95] Andy Gill and Simon Marlow. “Happy: the parser generator for Haskell”. In: *University of Glasgow* (1995).
- [Hol14] Timo von Holtz. *llvm-general-quote*. 2014. URL: <https://github.com/tvh/language-llvm-quote>.
- [Lat02] Chris Arthur Lattner. “LLVM: An infrastructure for multi-stage optimization”. MA thesis. University of Illinois, 2002. URL: <https://llvm.org/svn/llvm-project/www-pubs/trunk/2002-12-LattnerMSThesis-book.pdf>.
- [Mai07] Geoffrey Mainland. “Why it’s nice to be quoted: quasiquoting for haskell”. In: *Proceedings of the ACM SIGPLAN workshop on Haskell workshop*. ACM. 2007, pp. 73–82. URL: <https://www.cs.drexel.edu/~mainland/publications/mainland07quasiquoting.pdf>.
- [MC] Geoffrey Mainland and Manuel MT Chakravarty. *language-c-quote*. URL: <http://hackage.haskell.org/package/language-c-quote>.
- [RWZ88] Barry K Rosen, Mark N Wegman, and F Kenneth Zadeck. “Global value numbers and redundant computations”. In: *Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM. 1988, pp. 12–27. URL: <http://www.cs.wustl.edu/~cytron/cs531/Resources/Papers/valnum.pdf>.
- [Sca13] Benjamin Scarlet. *llvm-general*. 2013. URL: <http://hackage.haskell.org/package/llvm-general>.