

Master Thesis

An LLVM Backend for Accelerate



Christian-Albrechts-Universität zu Kiel
Department of Computer Science
Programming Languages and Compiler Construction

student: **Timo von Holtz**
advised by: Priv.-Doz. Dr. Frank Huch
Assoc. Prof. Dr. Manuel M. T. Chakravarty

Kiel, July 2, 2014

Selbstständigkeitserklärung

Ich erkläre hiermit, dass ich die vorliegende Arbeit selbstständig und nur unter Verwendung der angegebenen Literatur und Hilfsmittel angefertigt habe.

July 2, 2014

Timo von Holtz

Todo list

Write Introduction	1
write about Accelerate	8
longer introduction to monadic code generation	9
write about llvm-general-typed	9
update figure to correct ordering of bbs	11

Contents

1. Introduction	1
2. Technologies	3
2.1. LLVM	3
2.1.1. LLVM IR	3
2.1.2. Vectorization	5
2.1.3. llvm-general	8
2.2. Accelerate	8
3. Contributions	9
3.1. accelerate	9
3.2. llvm-general	9
3.3. llvm-general-quote	9
4. Implementation	15
4.1. Quasiquoter	15
4.2. Extension for-loop	15
4.3. SSA	15
5. Skeletons	17
5.1. map	17
5.2. fold	17
5.3. scan	17
6. Conclusion	19
6.1. Related Work	19
A. Listings	21
References	23

List of Figures

2.1. sum as a C function	3
2.2. sum as a LLVM	4
2.3. Vectorized Sum	6
2.4. Vectorized Sum	7
2.5. Vectorized Sum	8
3.1. Monadic generation of for loop	10
3.2. For Loop using <i>llvm-general-quote</i>	11
3.3. Expanded For Loop	12
3.4. For Loop using <i>llvm-general-quote</i> and SSA	13
3.5. Expanded For Loop (SSA)	13

List of Tables

1. Introduction

Write Introduction

2. Technologies

2.1. LLVM

LLVM[Lat02] is a compiler infrastructure written in C++. In contrast to GCC it is designed to be used as a library by compilers. Originally implemented for C and C++, the language-agnostic design (and the success) of LLVM has since spawned a wide variety of front ends: languages with compilers that use LLVM include ActionScript, Ada, D, Fortran, OpenGL Shading Language, Haskell, Java bytecode, Julia, Objective-C, Python, Ruby, Rust, Scala and C.

2.1.1. LLVM IR

LLVM defines it's own language to represent programs. It uses Static Single Assignment (SSA) form.[AWZ88; RWZ88] A program is said to be in SSA form if each of its variables is defined exactly once, and each use of a variable is dominated by that variable's definition. SSA form greatly simplifies many dataflow optimizations because only a single definition can reach a particular use of a value, and finding that definition is trivial.

To get idea of how this looks like in practice, let's look at an example. Figure 2.1 shows a simple C function to sum up the elements of an array. The corresponding LLVM code is shown in figure 2.2.

The first obvious difference is the lack of sophisticated control structures. In LLVM every function is divided into basic blocks. A basic block is a continuous stream of instructions with a terminator at the end. Instructions are add, mult, call, ..., but also

```
1 int sum(int* a, int length) {  
2     int x = 0;  
3     for (int i=0;i<length;i++) {  
4         x += a[i];  
5     }  
6     return x;  
7 }
```

Figure 2.1.: sum as a C function

```

1  define i32 @sum(i32* nocapture readonly %a, i32 %length) {
2  entry:
3      %cmp4 = icmp sgt i32 %length, 0
4      br i1 %cmp4, label %for.body, label %for.end
5
6  for.body:                                ; preds = %entry, %for.body
7      %indvars.iv = phi i64 [ %indvars.iv.next, %for.body ], [ 0, %entry ]
8      %x.05 = phi i32 [ %add, %for.body ], [ 0, %entry ]
9      %arrayidx = getelementptr inbounds i32* %a, i64 %indvars.iv
10     %0 = load i32* %arrayidx, align 4
11     %add = add nsw i32 %0, %x.05
12     %indvars.iv.next = add nuw nsw i64 %indvars.iv, 1
13     %lftr.wideiv = trunc i64 %indvars.iv.next to i32
14     %exitcond = icmp eq i32 %lftr.wideiv, %length
15     br i1 %exitcond, label %for.end, label %for.body
16
17 for.end:                                ; preds = %for.body, %entry
18     %x.0.lcssa = phi i32 [ 0, %entry ], [ %add, %for.body ]
19     ret i32 %x.0.lcssa
20 }

```

Figure 2.2.: sum as a LLVM

call. Terminators can either be used to jump to another block (branch) or return to the calling function.

To allow for dynamic control flow, there are Φ -nodes. These specify the value of a new variable depending on what the last block was. The Φ -nodes in the SSA are represented with **phi**-instructions. In LLVM, these have to precede every other instruction in a given basic block.

Types

The LLVM type system is one of the most important features of the intermediate representation. Being typed enables a number of optimizations to be performed directly, without having to do extra analyses on the side before the transformation.

Important Types are:

- **void**, which represents no value
- integers with specified length N: **iN**
- floating point numbers: **half**, **float**, **double**, ...
- pointers: **<type> ***
- function types: **<returntype> (<parameter list>)**
- vector types: **< <# elements> x <elementtype> >**
- array types: **[<# elements> x <elementtype>]**
- structure types: **{ <type list> }**

2.1.2. Vectorization

Modern CPUs all have SIMD units to execute an instruction on multiple datasets in parallel. Using these units is easy with LLVM. All operations (**add**, **fadd**, **sub**, ...) can be used with vector arguments the same way as with scalar arguments.

To manually exploit this can be tricky however. LLVM has multiple strategies to fuse similar instructions or tight inner loops into vectorized code.[LLV14]

Loop Vectorizer

The Loop Vectorizer tries to vectorize tight inner loops. To get an idea of how these work, let's look at the example from figure 2.1. To build the sum, every element of the array is added to an accumulator. Since addition is associative and commutative, the additions can be reordered. Given a vector width of 2, the sum of $2 * n$ and $2 * n + 1$ are calculated in parallel and then added together. In addition to this, the vectorizer will also unroll the inner loop to make fewer jumps necessary. Figure 2.3 shows the corresponding LLVM code without loop unroll, as this increases code size dramatically.

Apart from reductions, LLVM can also vectorize the following:

- Loops with unknown trip count
- Runtime Checks of Pointers
- Reductions
- Inductions
- If Conversion
- Pointer Induction Variables
- Reverse Iterators
- Scatter / Gather
- Vectorization of Mixed Types
- Global Structures Alias Analysis
- Vectorization of function calls
- Partial unrolling during vectorization

A detailed description can be found at <http://llvm.org/docs/Vectorizers.html>

SLP Vectorizer

The goal of SLP vectorization (a.k.a. superword-level parallelism) is to combine similar independent instructions into vector instructions. Memory accesses, arithmetic operations, comparison operations, PHI-nodes, can all be vectorized using this technique. For example, the following function performs very similar operations on its inputs (a1, b1) and (a2, b2).

```

1  define i32 @sum(i32* nocapture readonly %a, i32 %length) #0 {
2  entry:
3      %cmp4 = icmp sgt i32 %length, 0
4      br i1 %cmp4, label %for.body.preheader, label %for.end
5
6  for.body.preheader:                                ; preds = %entry
7      %0 = add i32 %length, -1
8      %1 = zext i32 %0 to i64
9      %2 = add i64 %1, 1
10     %end.idx = add i64 %1, 1
11     %n.vec = and i64 %2, 8589934590
12     %cmp.zero = icmp eq i64 %n.vec, 0
13     br i1 %cmp.zero, label %middle.block, label %vector.body
14
15  vector.body:                                        ; preds = %
16     for.body.preheader, %vector.body
17     %index = phi i64 [ %index.next, %vector.body ], [ 0, %for.body.preheader ]
18     %vec.phi = phi <2 x i32> [ %5, %vector.body ], [ zeroinitializer, %
19         for.body.preheader ]
20     %3 = getelementptr inbounds i32* %a, i64 %index
21     %4 = bitcast i32* %3 to <2 x i32>*
22     %wide.load = load <2 x i32>* %4, align 4
23     %5 = add nsw <2 x i32> %wide.load, %vec.phi
24     %index.next = add i64 %index, 2
25     %6 = icmp eq i64 %index.next, %n.vec
26     br i1 %6, label %middle.block, label %vector.body
27
28  middle.block:                                        ; preds = %vector.body, %
29     for.body.preheader
30     %resume.val = phi i64 [ 0, %for.body.preheader ], [ %n.vec, %vector.body ]
31     %rdx.vec.exit.phi = phi <2 x i32> [ zeroinitializer, %for.body.preheader ],
32         [ %5, %vector.body ]
33     %rdx.shuf = shufflevector <2 x i32> %rdx.vec.exit.phi, <2 x i32> undef, <2
34         x i32> <i32 1, i32 undef>
35     %bin.rdx = add <2 x i32> %rdx.vec.exit.phi, %rdx.shuf
36     %7 = extractelement <2 x i32> %bin.rdx, i32 0
37     %cmp.n = icmp eq i64 %end.idx, %resume.val
38     br i1 %cmp.n, label %for.end, label %for.body
39
40  for.body:                                            ; preds = %middle.block, %
41     for.body
42     %indvars.iv = phi i64 [ %indvars.iv.next, %for.body ], [ %resume.val, %
43         middle.block ]
44     %x.05 = phi i32 [ %add, %for.body ], [ %7, %middle.block ]
45     %arrayidx = getelementptr inbounds i32* %a, i64 %indvars.iv
46     %8 = load i32* %arrayidx, align 4
47     %add = add nsw i32 %8, %x.05
48     %indvars.iv.next = add nuw nsw i64 %indvars.iv, 1
49     %lftr.wideiv = trunc i64 %indvars.iv.next to i32
50     %exitcond = icmp eq i32 %lftr.wideiv, %length
51     br i1 %exitcond, label %for.end, label %for.body
52
53  for.end:                                            ; preds = %for.body, %
54     middle.block, %entry
55     %x.0.lcssa = phi i32 [ 0, %entry ], [ %add, %for.body ], [ %7, %
56         middle.block ]
57     ret i32 %x.0.lcssa
58 }

```

Figure 2.3.: Vectorized Sum

```

1  define void @foo(double %a1, double %a2, double %b1, double %b2, double* %A)
    {
2  entry:
3      %add = fadd double %a1, %b1
4      %mul = fmul double %add, %a1
5      %div = fdiv double %mul, %b1
6      %mul1 = fmul double %b1, 5.000000e+01
7      %div2 = fdiv double %mul1, %a1
8      %add3 = fadd double %div, %div2
9      store double %add3, double* %A, align 8
10     %add4 = fadd double %a2, %b2
11     %mul5 = fmul double %add4, %a2
12     %div6 = fdiv double %mul5, %b2
13     %mul7 = fmul double %b2, 5.000000e+01
14     %div8 = fdiv double %mul7, %a2
15     %add9 = fadd double %div6, %div8
16     %arrayidx10 = getelementptr inbounds double* %A, i64 1
17     store double %add9, double* %arrayidx10, align 8
18     ret void
19 }

```

Figure 2.4.: Vectorized Sum

```

1  void foo(double a1, double a2, double b1, double b2, double *A) {
2      A[0] = a1*(a1 + b1)/b1 + 50*b1/a1;
3      A[1] = a2*(a2 + b2)/b2 + 50*b2/a2;
4  }

```

The SLP vectorizer may combine these into vector operations. Figures 2.4 and 2.5 show the the corresponding LLVM before and after SLP vectorization.

Fast-Math Flags

In order for the Loop vectorizer to work, the operations involved need to be associative. When dealing with floating point numbers, the basic operations like **fadd** and **fmul** don't satisfy this condition. To vectorize the code regardless, LLVM has the notion of fast-math-flags. These tell the optimizer to assume certain properties that aren't true in general.

Available flags are[LLV]:

- nnan No NaNs - Allow optimizations to assume the arguments and result are not NaN. Such optimizations are required to retain defined behavior over NaNs, but the value of the result is undefined.
- ninf No Infs - Allow optimizations to assume the arguments and result are not +/-Inf. Such optimizations are required to retain defined behavior over +/-Inf, but the value of the result is undefined.

```
1  define void @foo(double %a1, double %a2, double %b1, double %b2, double* %A)
2  {
3  entry:
4      %0 = insertelement <2 x double> undef, double %a1, i32 0
5      %1 = insertelement <2 x double> %0, double %a2, i32 1
6      %2 = insertelement <2 x double> undef, double %b1, i32 0
7      %3 = insertelement <2 x double> %2, double %b2, i32 1
8      %4 = fadd <2 x double> %1, %3
9      %5 = fmul <2 x double> %4, %1
10     %6 = fdiv <2 x double> %5, %3
11     %7 = fmul <2 x double> %3, <double 5.000000e+01, double 5.000000e+01>
12     %8 = fdiv <2 x double> %7, %1
13     %9 = fadd <2 x double> %6, %8
14     %10 = bitcast double* %A to <2 x double>*
15     store <2 x double> %9, <2 x double>* %10, align 8
16     ret void
17 }
```

Figure 2.5.: Vectorized Sum

nsz No Signed Zeros - Allow optimizations to treat the sign of a zero argument or result as insignificant.

arcp Allow Reciprocal - Allow optimizations to use the reciprocal of an argument rather than perform division.

fast Fast - Allow algebraically equivalent transformations that may dramatically change results in floating point (e.g. reassociate). This flag implies all the others.

2.1.3. llvm-general

There are multiple bindings to LLVM in Haskell. The most complete is `llvm-general`.^[Sca13] Instead of exposing the LLVM API directly, it uses an ADT to represent LLVM IR. This can then be translated into the corresponding C++ object. It also supports the other way, transforming the object back into the ADT.

Using an ADT instead of writing directly to the C++ object has many advantages. This way, the code can be produced and manipulated outside the IO context. It is also much easier to reason about within Haskell.

2.2. Accelerate

write about Accelerate

Similar to Repa^[Kel+10], uses gang workers.^[Cha+07]

3. Contributions

3.1. accelerate

- StorableArray (errors in caching)

3.2. llvm-general

- Targetmachine (Optimization)
- fast-math

3.3. llvm-general-quote

When writing a compiler using LLVM in Haskell there is a good tutorial on how to do it at <http://www.stephendiehl.com/llvm/>. It uses *llvm-general* to interface with LLVM. The Idea followed in this tutorial is to use a monadic generator to produce the AST.

Figure 3.1 shows how to implement a simple for loop using monadic generators.

longer introduction to monadic code generation

Unfortunately, this produces a lot of boilerplate code. We have to define the basic blocks manually and add the instructions one by one. This has some obvious drawbacks, as the code can get unreadable pretty quickly.

Another approach would be to use an EDSL.

write about llvm-general-typed

I propose a third approach using quasiquotation[Mai07]. The idea behind quasiquotation is, that you can define a DSL with arbitrary syntax, which you can then directly transform into Haskell data structures. This is done at compile-time, so you get the same type safety as writing the AST by hand.

Figure 3.3 shows how to implement a simple function in LLVM using quasiquotation. Compared to the other 2 solutions, this has already the advantage of being very close to the produced LLVM IR.

```

for :: Type                                -- type of the index
      → Operand                             -- starting index
      → (Operand → CodeGen Operand)         -- loop test to keep going
      → (Operand → CodeGen Operand)         -- increment the index
      → (Operand → CodeGen ())              -- body of the loop
      → CodeGen ()

for ti start test incr body = do
  loop ← newBlock "for.top"
  exit ← newBlock "for.exit"

  -- entry test
  c ← test start
  top ← cbr c loop exit

  -- Main loop
  setBlock loop
  c_i ← freshName
  let i = local c_i

  body i
  i' ← incr i
  c' ← test i'
  bot ← cbr c' loop exit
  – ← phi loop c_i ti [(i', bot), (start, top)]
  setBlock exit

```

Figure 3.1.: Monadic generation of for loop

```

1  define i32 @foo(i32 %x) #0 {
2  entry:
3    br label %for.cond
4
5  for.cond:                                ; preds = %for.inc, %entry
6    %res.0 = phi i32 [ 0, %entry ], [ %add, %for.inc ]
7    %i.0 = phi i32 [ 0, %entry ], [ %inc, %for.inc ]
8    %cmp = icmp slt i32 %i.0, %x
9    br i1 %cmp, label %for.body, label %for.end
10
11 for.body:                                ; preds = %for.cond
12   %add = add nsw i32 %res.0, %x
13   %inc = add nsw i32 %i.0, 1
14   br label %for.cond
15
16 for.end:                                ; preds = %for.cond
17   ret i32 %res.0
18 }

```

```

1  [llg|
2  define i64 @foo(i64 %start, i64 %end) {
3    entry:
4      br label %for
5
6    for:
7      for i64 %i in %start to %end with i64 [0,%entry] as %x {
8        %y = add i64 %i, %x
9        ret i64 %y
10     }
11
12    exit:
13      ret i64 %y
14  }
15  |]
```

Figure 3.2.: For Loop using *llvm-general-quote*

Without the ability to reference Haskell variables, this would be fairly useless in most cases. But quasiquotation allows for antiquotation as well. This means you can still reference arbitrary Haskell variables from within the quotation. Using this the following are equivalent:

- `[llinstr| add i64 %x, 1 |]`
- `let y = 1 in [llinstr| add i64 %x, $opr:y |]`

3.3.1. control-structures

You still have to specify the Φ -nodes by hand. This is fairly straight forward in a simple example, but can get more complicated very quickly.

The real goal is a language that “feels” like a high-level language, but can be trivially translated into LLVM. This means

- using LLVM instructions unmodified.
- introduce higher-level control structures like for, while and if-then-else.

Figure 3.2 shows a for loop using this approach. The loop-variable (%x in this case) is specified explicitly and can be referenced inside and after the for-loop. To update the %x, I overload the return statement to specify which value should be propagated. At the end of the for-loop the code automatically jumps to the next block. Figure 3.3 shows the produced LLVM code. This is clearly more readable.

```
1 define i64 @foo(i64 %start, i64 %end) {  
2   entry:  
3     br label %for  
4  
5   for:                                     ; preds = %for.body, %entry  
6     %i = phi i64 [ %i.new, %for.body ], [ %start, %entry ]  
7     %x = phi i64 [ %y, %for.body ], [ 0, %entry ]  
8     %for.cond = icmp ule i64 %i, %end  
9     %i.new = add nuw nsw i64 %i, 1  
10    br i1 %for.cond, label %for.body, label %for.end  
11  
12   for.body:                               ; preds = %for  
13     %y = add i64 %i, %x  
14     br label %for  
15  
16   for.end:                               ; preds = %for  
17     ret i64 %x  
18 }
```

Figure 3.3.: Expanded For Loop

3.3.2. SSA

The above approach unfortunately only works for some well-defined cases. With multiple loop-variables for example, it quickly becomes cluttered. On top of this, it relies on some “magic” to resolve the translation and legibility suffers.

The main reason why it is not possible to define a better syntax is that LLVM code has to be in SSA form. This means, that

1. every variable name can only be written to once.
2. Φ -nodes are necessary where control flow merges.

To loosen this restriction, I implemented SSA recovery pass as part of the quasiquote. This means that I can now reassign variables inside the LLVM code. Using this, I was able to define a much more simple syntax for the `for`-loop. The placement of Φ -nodes is not optimal, but redundant ones can be easily remove by LLVM’s `InstCombine` pass. Figure 3.4 and 3.5 show the quoted and the produced code after `InstCombine` respectively.

3.3.3. syntax extensions

`llvm-general-quote` supports all `llvm` instructions. I added the following to the syntax:

- direct assignment: `<var> = <ty> <val>`
- if: `if <cond> { <instructions> } [else { <instructions> }]`
- while: `while <cond> { <instructions> }`


```

1  [llg]
2  define i64 @foo(i64 %start, i64 %end) {
3      entry:
4          %x = i64 0
5
6      for:
7          for i64 %i in %start to %end {
8              %x = add i64 %i, %x
9          }
10
11     exit:
12         ret i64 %x
13 }
14 []

```

Figure 3.4.: For Loop using *llvm-general-quote* and SSA

```

1  define i64 @foo(i64 %start, i64 %end) {
2      entry:
3          br label %for.head
4
5      for.head:                                     ; preds = %n0, %entry
6          %x.12 = phi i64 [ 0, %entry ], [ %x.6, %n0 ]
7          %i.4 = phi i64 [ %start, %entry ], [ %i.9, %n0 ]
8          %for.cond.3 = icmp slt i64 %i.4, %end
9          br i1 %for.cond.3, label %n0, label %for.end
10
11     n0:                                           ; preds = %for.head
12         %x.6 = add i64 %i.4, %x.12
13         %i.9 = add nuw nsw i64 %i.4, 1
14         br label %for.head
15
16     for.end:                                     ; preds = %for.head
17         ret i64 %x.12
18 }

```

Figure 3.5.: Expanded For Loop (SSA)

3. Contributions

- for: for <ty> <var> in <val1> (**to** | downto)<val2> { <instructions> }

4. Implementation

4.1. Quasiquoter

The design of *llvm-general-quote* is inspired by *language-c-quote*, which is also used in the cuda implementation of Accelerate. I use “Happy” and *Alex*.

4.2. Extension for-loop

4.3. SSA

The standard method to achieve this is to use stack allocated variables instead of registers. The LLVM optimizer then uses the mem2reg pass to transform these into registers, adding the necessary phi-nodes in the process. [Bra+13]

5. Skeletons

5.1. map

5.2. fold

5.3. scan

[LF80] My plan is the following:

```
1 void scan(double* in, double* out, double* tmp, unsigned length, unsigned
  start, unsigned end, unsigned tid) {
2   double acc = 0;
3   if (end==length) {
4     tmp[0] = 0;
5   } else {
6     for (unsigned i=start; i<end; i++) {
7       acc += in[i];
8     }
9     tmp[tid+1] = acc;
10  }
11  printf("block");
12  acc = tmp[tid];
13  for (unsigned j=start; j<end; j++) {
14    acc += in[j];
15    out[j] = acc;
16  }
17 }
18
19 void scanAlt(double* in, double* out, unsigned length, unsigned start,
  unsigned end) {
20   double acc = 0;
21   for (unsigned j=start; j<end; j++) {
22     acc += in[j];
23     out[j] = acc;
24   }
25   printf("block");
26   double add=0;
27   if (start>0) {
28     add = out[start-1];
29   }
30   for (unsigned i=start; i<end; i++) {
31     out[i] += add;
```

```
32     }  
33 }
```

6. Conclusion

6.1. Related Work

A. Listings

References

- [AWZ88] ALPERN, Bowen ; WEGMAN, Mark N ; ZADECK, F Kenneth: Detecting equality of variables in programs. In: *Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM. 1988, pp. 1–11. URL: <https://courses.cs.washington.edu/courses/cse501/04wi/papers/alpern-popl88.pdf>
- [Bra+13] BRAUN, Matthias et al.: Simple and Efficient Construction of Static Single Assignment Form. In: *Proceedings of the International Conference on Compiler Construction*. Lecture Notes in Computer Science. 2013. URL: <http://pp.info.uni-karlsruhe.de/uploads/publikationen/braun13cc.pdf>
- [Cha+07] CHAKRAVARTY, Manuel MT et al.: Data Parallel Haskell: a status report. In: *Proceedings of the 2007 workshop on Declarative aspects of multicore programming*. ACM. 2007, pp. 10–18. URL: <http://www.cse.unsw.edu.au/~chak/papers/data-parallel-haskell.pdf>
- [Die14] DIEHL, Stephen: *Implementing a JIT Compiled Language with Haskell and LLVM*. Jan. 2014. URL: <http://www.stephendiehl.com/llvm/>
- [DM] DORNAN, Chris ; MARLOW, Simon: *Alex: A lexical analyser generator for Haskell*. URL: <http://www.haskell.org/alex/>
- [GM95] GILL, Andy ; MARLOW, Simon: Happy: the parser generator for Haskell. In: *University of Glasgow* (1995)
- [Hol14] HOLTZ, Timo von: *llvm-general-quote*. 2014. URL: <https://github.com/tvh/language-llvm-quote>
- [Kel+10] KELLER, Gabriele et al.: Regular, shape-polymorphic, parallel arrays in Haskell. In: *ACM Sigplan Notices*. Vol. 45. 9. ACM. 2010, pp. 261–272. URL: <http://www.cse.unsw.edu.au/~chak/papers/repaper.pdf>
- [LF80] LADNER, Richard E ; FISCHER, Michael J: Parallel prefix computation. In: *Journal of the ACM (JACM)* 27 (1980) Nr. 4, pp. 831–838. URL:

- <https://engr.smu.edu/~seidel/courses/cse8351/papers/LadnerFischer80.pdf>
- [Lat02] LATTNER, Chris Arthur: *LLVM: An infrastructure for multi-stage optimization*. MA thesis. University of Illinois, 2002. URL: <https://llvm.org/svn/llvm-project/www-pubs/trunk/2002-12-LattnerMSThesis-book.pdf>
- [LLV14] LLVM PROJECT: *Auto-Vectorization in LLVM*. [Online; accessed 1-July-2014]. 2014. URL: <http://llvm.org/docs/Vectorizers.html>
- [LLV] LLVM PROJECT: *LLVM Language Reference Manual*. [Online; accessed 1-July-2014]. URL: <http://llvm.org/docs/LangRef.html>
- [Mai07] MAINLAND, Geoffrey: Why it's nice to be quoted: quasiquoting for haskell. In: *Proceedings of the ACM SIGPLAN workshop on Haskell workshop*. ACM. 2007, pp. 73-82. URL: <https://www.cs.drexel.edu/~mainland/publications/mainland07quasiquoting.pdf>
- [MC] MAINLAND, Geoffrey ; CHAKRAVARTY, Manuel MT: *language-c-quote*. URL: <http://hackage.haskell.org/package/language-c-quote>
- [RWZ88] ROSEN, Barry K ; WEGMAN, Mark N ; ZADECK, F Kenneth: Global value numbers and redundant computations. In: *Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM. 1988, pp. 12-27. URL: <http://www.cs.wustl.edu/~cytron/cs531/Resources/Papers/valnum.pdf>
- [Sca13] SCARLET, Benjamin: *llvm-general*. 2013. URL: <http://hackage.haskell.org/package/llvm-general>